

Monte Carlo Particle Lists : MCPL

T Kittelmann^{a,*}, E Klinkby^b, E B Knudsen^c, P Willendrup^c, X X Cai^{a,b},
K Kanaki^a

^a*European Spallation Source ERIC, Sweden*

^b*DTU Nutech, Technical University of Denmark, Denmark*

^c*DTU PHYSICS, Technical University of Denmark, Denmark*

Abstract

A binary format with lists of particle state information, for interchanging particles between various Monte Carlo simulation applications, is presented. Portable C code for file manipulation is made available to the scientific community, along with converters and plugins for several popular simulation packages.

PROGRAM SUMMARY

Manuscript Title: Monte Carlo Particle Lists : MCPL

Authors: T Kittelmann, E Klinkby, E B Knudsen, P Willendrup, X X Cai and K Kanaki

Program Title: MCPL

Journal Reference:

Catalogue identifier:

Licensing provisions: CC0 for core MCPL, see LICENSE file for details.

Programming language: C and C++

Operating system: Linux, OSX, Windows

Keywords:

MCPL, Monte Carlo, particles, data storage, simulation, interchange format, C++, C, Geant4, MCNP, McStas, McXtrace

Classification: 4, 9, 11, 17

External routines/libraries: Geant4, MCNP, McStas, McXtrace

Nature of problem:

*Corresponding author. *Email address:* thomas.kittelmann@esss.se

Persistification of particle states in Monte Carlo simulations, for interchange between simulation packages or for reuse within a single package.

Solution method:

Binary interchange format with associated code written in portable C along with tools and interfaces for relevant simulation packages.

1. Introduction

The usage of Monte Carlo simulations, to describe the transport and interaction of particles and radiation, is a powerful and popular technique, finding use throughout a wide range of fields – including but not limited to both high energy and nuclear physics as well as space and medical sciences^[*todo: missing text or text in need of edit here*]. Naturally, a plethora of different frameworks and applications exists for carrying out these simulations (cf. section 3 for examples), with implementations in different languages and domains ranging from general purpose to highly specialised field- and application-specific.

A common principle used in the implementation of these applications is the representation of particles by a set of state parameters – usually including at least particle type, time coordinate, position and momentum vectors – and a suitable representation of the geometry of the problem (either via descriptions of actual surfaces and volumes in a virtual three-dimensional space, or through suitable parameterisations). In the simplest scenario where no variance-reduction techniques are employed, simulations are typically carried out by proceeding iteratively in steps from an initial set of particles states, with the state information being updated along the way as a result of the pseudo-random or deterministic modelling of processes affecting the particle. The modelling can represent particle self-interactions, interactions with the material of the simulated geometry, or simply its forward transport through the geometry, using either simple ray-tracing techniques or more complicated trajectory calculations as appropriate. In addition to a simple update of state parameters, the modelling can result in termination of the simulation for the given particle or in the

25 creation of new secondary particle states, which will in turn undergo simulation themselves.

Occasionally, use-cases arise in which it would be beneficial to be able to capture a certain subset of particle states present in a given simulation, in order to continue their simulation at a later point in either the same or a different
30 framework. Such capabilities have typically been implemented using custom application-specific means of data exchange, often involving the tedious writing of custom input and output hooks for the specific frameworks and use-cases in question. Here is instead presented a standard format for exchange of particle state data, *Monte Carlo Particle Lists* (MCPL), which is intended to replace the
35 plethora of custom converters with a more convenient scenario in which experts of each framework implement converters to the common format, as a one-time effort. The idea being that users of the various frameworks then gain the ability to simply activate those preexisting and validated converters in order to carry out their work.

40 The present work originated in the needs for simulations at neutron scattering facilities, where a multitude of simulation frameworks are typically used to describe the various components from neutron production to detection, but historically other conceptually similar formats have been and is used in high energy physics to communicate particle states between event generators and
45 detector simulations [1, 2, 3]. Here, the formats were mainly developed for in-process communication and also had to accommodate the potential description of intermediary unphysical or bound particles, as well as potentially keeping meta-data concerning the simulated history of a given particle in the record. For these reasons, for reasons of portability and ease of integration, and due to
50 the lack of a common well-defined on-disk format, these existing solutions were deemed unfit for the goals of the work presented here: A compact yet flexible on-disk binary format for particle state information, able to accommodate a wide range of use-cases with close to optimal storage requirements. The accompanying code with which to access and manipulate it, should be small, efficient
55 and easily integrated into existing codes and build systems. Consequently, it

was chosen to implement the format through a set of C functions declared in a single header file, `mcpl.h`, and implemented in a single file, `mcpl.c`.

2. The MCPL format

MCPL is a binary file format, in which a header section with configuration and meta-data is followed by a data section where the state information of the contained particles is kept. Data compression is available but optional (cf. section 2.4). The uncompressed storage size of a particle entry in the data section is determined by global settings in the header section, and depends on what exact information is stored for the particles in a given file, as will be discussed shortly. Within a given file, all particle entries will always be of equal length, allowing for trivial calculation of the absolute data location for a particle at a given index in the file – and thus for efficient seeking and skipping between particles if desired. As it is expected that MCPL files will always be read by calls to the functions in `mcpl.h`, no attempt will be made here to provide a complete specification of the binary layout of data in the files. Instead, interested readers are referred to the implementation of the functions `mcpl_write_header` and `mcpl_add_particle` in `mcpl.c`.

2.1. Information available

The information available in the file header is indicated in Table 1: A unique 4-byte magic number identifying the format always starts all files, and is followed by the format version, the endianness (*little* or *big*) in which numbers in the file are stored, and the number of particles in the file. The versioning provides a clear path for future updates to the format, without losing the ability to read files created with previous versions of the MCPL code, and the endianness information prevents interpretation errors on different machines (although at present, most consumer platforms are little-endian).¹ Next comes four options

¹In the current implementation, reading a little-endian MCPL file on a big-endian machine or vice versa triggers an error message. It is envisioned that a future version of the MCPL code

File header information	
<i>Field</i>	<i>Description</i>
File type magic number 0x4d43504c ("MCPL")	All MCPL files start with this 4-byte word.
Version	File format version.
Endianness	Whether numbers in file are in little- or big-endian format.
Number of particles in file	64 bit integer.
Flag : Particles have polarisation info	If false, all loaded particles will have polarisation vectors (0,0,0).
Flag : Particles have "userflags" field	If false, all loaded particles will have userflags 0x00000000.
Flag : Particle info use double-precision	If true, floating points storage use double-precision.
Global pdgcode	If this 32 bit integer is non-zero, all loaded particles will have this pdgcode.
Source name	String indicating the application which created the MCPL file.
Comments	A variable number of comments (strings) added at file creation.
Binary blobs	A variable number of binary data blobs, indexed by keys (strings). This allows arbitrary custom data to be embedded.

Table 1: Information available in the header section of MCPL files.

affecting what data is stored per-particle, which will be discussed in the next paragraph. Finally, the header contains several options for embedding custom free-form information: First of all, the source name, in the form of a single string
85 containing the name and perhaps version of the application which created the file. Secondly, any number of strings can be added as human readable comments, and, thirdly, any number of binary data blobs can be added, each identified by a string key. The MCPL format itself provides no restrictions on what data, if any, can be stored in these binary blobs, but useful content could for instance
90 be a copy of configuration data used by the source application when the given file was produced, kept for later reference.

Table 2 shows the state information available per-particle in MCPL files, along with the storage used by each field. Particle position, direction, kinetic energy, time and weight are always stored². Polarisation vectors and so-called *user-flags*
95 in the form of 32 bit integers are only stored when relevant flags in the header are

could instead transparently correct the endianness at load time.

²Note that a valid alternative to storing the directional unit vector along with the kinetic energy would have been the momentum vector. However, our choice is consistent with the variables used in interfaces of both MCNP and Geant4, and means that the MCPL-SSW converter discussed in section 3.2 can be implemented without access to an unwieldy database of particle and isotope masses.

Particle state information		
<i>Field</i>	<i>Description</i>	<i>Bytes of storage used per entry (FP = 4 or 8 bytes)</i>
PDG code	32 bit integer indicating particle type.	0 or 4
Position	Vector, values in centimetres.	3FP
Direction	Unit vector along the particle momentum.	2FP
Kinetic energy	Value in MeV.	1FP
Time	Value in milliseconds.	1FP
Weight	Weight or intensity.	1FP
Polarisation	Vector.	0 or 3FP
User flags	32 bit integer with custom info.	0 or 4

Table 2: Particle state information available and uncompressed storage requirements for each entry in the data section of MCPL files.

enabled. Likewise, the particle type information in the form of so-called PDG codes are only stored explicitly in each entry when a global PDG code was not specified in the header. The PDG codes must follow the scheme developed by the Particle Data Group in [4, ch. 42], which is inarguably the most comprehensive and widely adopted standard for particle type encoding in simulations. Finally, again depending on a flag in the header, particle information use either single (4 bytes) or double precision (8 bytes) storage for floating point numbers. All in all, summing up the numbers in the last column of Table 2, particles are seen to consume between 32 and 96 bytes of uncompressed storage space per entry. The MCPL format is thus designed to be flexible enough to handle use-cases requiring a high level of detail in the particle state information, without imposing excessive storage requirements on less demanding scenarios.

Packing of the three-dimensional unit directional vector into just two floating point numbers of storage is carried out via a custom octahedral packing algorithm inspired by [5, 6]. This packing avoids the very significant loss of numerical precision otherwise resulting when using the straight-forward and widespread solution of storing two Cartesian components, u_x and u_y , directly and recovering the magnitude of the third by the expression $|u_z| = \sqrt{1 - u_x^2 - u_y^2}$.

2.2. Accessing or creating MCPL files programmatically

115 The main feature provided by the implementation of MCPL in `mcpl.h` and `mcpl.c` is naturally the ability to create new MCPL files and access the contents of existing ones, using a set of dedicated functions. No matter which settings were chosen when a given MCPL file was created, the interface for accessing the header and particle state information within it is the same, as can be seen in

120 Listing 1: After obtaining a file handle via `mcpl_open_file`, a pointer to an `mcpl_particle_t` struct, whose fields contain the state information available for a given particle, is returned by calling `mcpl_read`. This also advances the position in the file, and returns a null-pointer when there are no more particles in the file, ending the loop. If a file was created with either polarisation vectors or

125 user-flags disabled, the corresponding fields on the particle will contain zeroes (thus representing polarisation info with null-vectors and user-flags with an integer with no bits enabled). All floating point fields on `mcpl_particle_t` are represented with a double-precision type, but the actual precision of the numbers will obviously be limited to that used in the input file. In addition

130 to the interface illustrated by Listing 1, functions can be found in `mcpl.h` for accessing any information available in the file header (see Table 1), or for seeking and skipping to particles at specific positions in the file, rather than simply iterating through the full file.

Code creating MCPL files is typically slightly more involved, as the creation

135 process also involves deciding on the values of the various header flags and filling of free-form fields like source name and comments. An example producing a file with 1000 particles is shown in Listing 2. The most important part of the procedure is to first obtain a file handle through a call to `mcpl_create_file`, configure the header and overall flags, and prepare a zero-initialised instance

140 of `mcpl_particle_t`. Next comes the loop filling the particles into the file, which happens by updating the numbers on the `mcpl_particle_t` instance as needed, and passing it to `mcpl_add_particle` each time. At the end, a call to `mcpl_close_outfile` finishes up by flushing all internal buffers to disk and updating the field containing the number of particles at the beginning of the

Listing 1: Simple example for looping over all particles in an existing MCPL file

```
#include "mcpl.h"

void example()
{
    mcpl_file_t f = mcpl_open_file("myfile.mcpl");

    const mcpl_particle_t* p;

    while ( ( p = mcpl_read(f) ) ) {

        /* Particle properties can here be accessed
           through the pointer "p":

           p->pdgcode
           p->position[k] (k=0,1,2)
           p->direction[k] (k=0,1,2)
           p->polarisation[k] (k=0,1,2)
           p->ekin
           p->time
           p->weight
           p->userflags
        */

    }

    mcpl_close_file(f);
}
```

145 file.

Should the unfortunate happen, that the program is aborted before the call to `mcpl_close_outfile`, any particles already written into the output file are normally recoverable: Upon opening such an incomplete file, the MCPL code detects that the actual size of the file is inconsistent with the value of field in the header containing the number of particles. Thus, it emits a warning message and
150 calculates a more appropriate value for the field, ignoring any partially written particle state entry at the end of the file. This ability to transparently correct incomplete files upon load also means that it is possible to inspect (with the `mcpltool` command discussed in section 2.3) or analyse files that are still being
155 created. To avoid seeing a warning each time a file left over from an aborted job is opened, `mcpl.h` also provides the function `mcpl_repair` which can be used to permanently correct the header of the file.

Likewise, `mcpl.h` also provides the function `mcpl_merge` which can be used to merge two compatible MCPL files into one, which might typically be useful

Listing 2: Simple example for creating an MCPL file with 1000 particles.

```
#include "mcpl.h"
#include <stdlib.h>

void example()
{
    mcpl_outfile_t f = mcpl_create_outfile("myfile.mcpl");
    mcpl_hdr_set_srcname(f, "MyAppName-1.0");

    /* Tune file options or add custom comments or
       binary data into the header:

        mcpl_enable_universal_pdgcode(f, myglobalpdgcode);
        mcpl_enable_userflags(f);
        mcpl_enable_polarisation(f);
        mcpl_enable_doubleprec(f);
        mcpl_hdr_add_comment(f, "Some comment.");
        mcpl_hdr_add_data(f, "mydatakey",
                        my_datalength, my_databuf)
    */

    mcpl_particle_t * p;
    p = (mcpl_particle_t*)calloc(sizeof(mcpl_particle_t), 1);

    int i;
    for ( i = 0; i < 1000; ++i ) {

        /* The following particle properties must
           always be set here:

            p->position[k] (k=0,1,2)
            p->direction[k] (k=0,1,2)
            p->ekin
            p->time
            p->weight

           These should also be set when required by
           file options:

            p->pdgcode
            p->userflags
            p->polarisation[k] (k=0,1,2)
        */

        mcpl_add_particle(f, p);
    }

    mcpl_close_outfile(f);
    free(p);
}
```

160 when gathering up the output of simulations carried out via parallel processing techniques. Compatibility here means that the files must have essentially identical header sections, except for the field holding the number of particles.

2.3. Accessing MCPL files from the command line

Compared with simpler text-based formats (e.g. ASCII files with data formatted in columns), one potential disadvantage of a binary data format like MCPL is the lack of an easy way to quickly inspect a file to investigate its contents. To alleviate this, the `mcpl.h` and `mcpl.c` files implement a function `int mcpl_tool(int argc, char** argv)` which, in a straight-forward manner, can be used to build a generic `mcpltool` command-line executable. Simply running this command on an MCPL file without specifying other arguments, results in a short summary of the file content being printed to standard output. This includes a listing of the first 10 contained particles, and an example of such a summary is provided in Listing 3: It is clear from the displayed metadata that the particles in the given file represents a transmission spectrum resulting from illumination of a block of lead with a 10 GeV proton beam in a **Geant4** [7, 8] simulation. The displayed header info and data columns should be mostly self-explanatory, noting that `(x,y,z)` indicates the particle position, `(ux,uy,uz)` its direction, and that the `pdgcode` column indeed show particle types typical in a hadronic shower: π^+ (211), γ (22), protons (2212), π^- (-211) and neutrons (2112). If the file had user-flags or polarisation vectors enabled, appropriate columns for those would be shown as well. Finally, note that the `36 bytes/particle` refers to uncompressed storage, and that in this particular case the file actually has a compression ratio of approximately 70%, meaning that about 25 bytes of on-disk storage is used per particle (cf. section 2.4).

Listing 3: Example output of running mcpltool, with no arguments on a specific MCPL file.[\[todo: missing text or text in need of edit here\]](#)

Opened MCPL file myoutput.mcpl.gz:

Basic info										
Format	: MCPL-2									
No. of particles	: 1106933									
Header storage	: 140 bytes									
Data storage	: 39848588 bytes									
Custom meta data										
Source	: "G4MCPLWriter [G4MCPLWriter]"									
Number of comments	: 1									
Number of blobs	: 0									
Particle data format										
Storage	: no									
User flags	: no									
Polarisation info	: no									
Fixed part. type	: no									
FP precision	: single									
Endianness	: little									
Storage	: 36 bytes/particle									
index	pdgcode	ekin[MeV]	x[cm]	y[cm]	z[cm]	ux	uy	uz	time[ns]	weight
0	211	487.02	-0.5898	1.835	20	-0.092407	0.20491	0.97441	7.3346e-07	1
1	22	1.5326	-0.18635	11.351	20	0.080441	0.56026	0.74672	1.0882e-06	1
2	22	3.9526	-0.43907	8.6473	20	-0.56616	0.50558	0.65104	1.0286e-06	1
3	22	0.82591	1.7444	9.7622	20	0.092099	0.79597	0.59829	1.0378e-06	1
4	22	1.1958	2.1806	8.6416	20	0.21997	0.66435	0.71432	1.0124e-06	1
5	22	1.2525	3.0949	7.7366	20	0.48903	0.30789	0.81512	1.013e-06	1
6	22	2.6247	3.948	5.681	20	0.62503	0.64221	0.44374	9.1152e-07	1
7	2212	824.28	-1.8797	-2.5124	20	-0.3077	-0.40486	0.861	7.6539e-07	1
8	-211	3459.8	-0.79521	0.91481	20	-0.13441	0.14438	0.98035	7.0618e-07	1
9	2112	0.30553	54.471	33.386	20	0.4862	0.011958	0.87377	0.00016442	1

185 By providing suitable arguments (found by the command `mcpltool --help`)
to `mcpltool`, it is possible to modify what information from the file is displayed.
This includes the possibility to change what particles from the file, if any, should
be listed, as well as the option to extract the contents of a given binary data
blob to standard output. The latter might be particularly handy when entire
190 configuration files have been embedded (cf. sections 3.2 and 3.3). Finally,
the `mcpltool` command also allows file merging and repairing, as discussed in
section 2.2.

Advanced functionality such as graphics display and interactive GUI-based
investigation or manipulation of the contents of MCPL files is not provided by
195 the `mcpltool`, since those would imply additional unwanted dependencies to
the core MCPL distribution, which is required by design to be light-weight and
widely portable. However, it is the hope that the existence of a standard format
like MCPL will encourage development of such tools, and indeed some already
exists in the in-house framework of the ESS Detector Group [9]. It is intended
200 for a future distribution of MCPL to include relevant parts of these tools as a
separate and optional component.

2.4. Compression

The utilisation of data compression in a format like MCPL is potentially an
important feature, since on-disk storage size could be a concern for some appli-
205 cations. Aiming to maximise flexibility, transparency and portability, optional
compression of MCPL files is simply provided by allowing whole-file compression
into the widespread GZIP format[10] (changing the file extension from `.mcpl` to
`.mcpl.gz` in the process). This utilises the DEFLATE compression algorithm[11]
which offers a good performance compromise with a reasonable compression
210 ratio and an excellent speed of compression and decompression.

Relying on a standard format such as GZIP means that, if needed, users can
avail themselves of existing tools (like the `gzip` and `gunzip` commands available
on most UNIX platforms) to change the compression state of an existing MCPL file.
However, when the code in `mcpl.c` is linked with the ubiquitous ZLIB[12, 13]

215 (cf. section 2.5), compressed MCPL files can be read directly. Additionally, for
convenience `mcpl.h` provides a function `mcpl_closeandgzip_outfile`, which
can be used instead of `mcpl_close_outfile` (cf. Listing 2) to ensure that newly
created MCPL files are automatically compressed if possible (either through a
call to an external `gzip` command or through custom ZLIB-dependent code,
220 depending on availability).

2.5. Build and deployment

It is the hope that eventually MCPL capabilities will be included upstream
in many applications, and that users of those consequently won't have to do
anything extra to start using it. As will be discussed in section 3, this is at
225 present the case for users of recent versions of `McStas`, `McXtrace` and the in-
house `Geant4`-based framework of the ESS Detector Group [9].

By design, it is expected that most developers wishing to add MCPL support
to their application will simply place copies of `mcpl.h` and `mcpl.c` into their
existing build system and include `mcpl.h` from either C or C++ code³. In order to
230 make the resulting binary code able to manipulate compressed files directly (cf.
section 2.4), the code in `mcpl.c` must usually be compiled against and linked
with an installation of ZLIB (see detailed instructions regarding build flags at
the top of `mcpl.c`). Alternatively, the MCPL distribution presented here contains
a “fat” auto-generated drop-in replacement for `mcpl.c` named `mcpl_fat.c`, in
235 which the source code of ZLIB has been included in its entirety⁴. Using this
somewhat larger file enables ZLIB-dependant code in MCPL even in situation
where ZLIB might not be otherwise available.

In addition to the core MCPL code, the MCPL distribution also contains a

³Compilation of `mcpl.c` can happen with any of the following standards: C99, C11, C++98, C++11, C++14, or later. In addition to those, `mcpl.h` is also C89 compatible. Note that on platforms where the standard C math function `sqrt` is provided in a separate library, that library must be available at link-time.

⁴Note that all ZLIB symbols have been prefixed, to guard against potential run-time clashes where a separate ZLIB is nonetheless loaded.

small file providing the `mcpltool` executable, C++ files implementing the `Geant4`
240 classes discussed in section 3.1, C files for the `mcpl2ssw` and `ssw2mcpl` executables discussed in section 3.2, and a few examples show-casing how user code might look.

Building of all of these parts should be straight-forward using standard tools, but a configuration file for `CMake`[14] which builds and installs everything is
245 nonetheless provided for reference and convenience.

Finally, “fat” single-file versions of all command line utilities (`mcpltool`, `mcpl2ssw` and `ssw2mcpl`) are also provided, containing both MCPL and ZLIB code within as appropriate. Thus, any of these single-file versions can be compiled directly into the corresponding command line executable, without any other
250 dependencies than a C compiler.

3. Application-specific converters and plugins

While the examples in Section 2.2 show how it is possible to manipulate MCPL files directly from code build with C or C++, it is not envisioned that most users will have to write such code themselves. Rather, in addition to using commonly
255 available tools (cf. section 2.3) to investigate the contents of files as needed, users would ideally simply use preexisting plugins and converters written by application-specific experts, to load particles from MCPL files into their given applications, or extract particles from those and into MCPL files. At the time of this initial public release of MCPL, four such applications are already MCPL-
260 aware in this manner: `Geant4`, `MCNP`, `McStas` and `McXtrace`[todo: missing text or text in need of edit here], and the details of the corresponding converters and plugins are discussed in the following sub-sections, after a few general pieces of advice for other implementers in the next paragraphs.

In order for MCPL files to be as widely exchangeable as possible, code loading
265 particles from MCPL files into a given Monte Carlo application should preferably be as accepting as possible. In particular, this means that warnings rather than errors should result if the input file contains PDG codes corresponding to

particle types that can not be handled by the application in question. As an example, a detailed MCNP or **Geant4** simulation of a moderated neutron source will typically produce files containing not only neutrons, but also gammas and other particles. It should certainly be possible to load such a file into a neutron-only simulation application like **McStas**, resulting in simulation of the contained neutrons (preferably with a warning or informative message drawing attention to some particles being ignored).

Applications employing parallel processing techniques, must pay particular attention when implementing file-based I/O, and this is naturally also the case when creating MCPL-aware plugins for them. However, the available functionality for merging of MCPL files makes the scenario of file creation particularly simple to implement: Each sub-task can simply write its own file, with the subsequent merging into a single file taking place during post-processing. For reading of particles in existing MCPL files, it is recommended that each sub-task performs a separate call `mcpl_open_file`, and use the skipping and seeking functionality to load just a subset of the particles within, as required. In the case of a multi-threading application, it is of course also possible to handle concurrent input or output directly through a single file handle. In this case, however, calls to `mcpl_add_particle` and `mcpl_read` must be protected against concurrent invocations with a suitable lock or mutex.

The following three sub-sections are dedicated to discussions of MCPL interfaces for specific Monte Carlo applications. The discussions will in each case presuppose familiarity with the application in question.

3.1. *Geant4* interface

In the most typical mode of working with the **Geant4** toolkit, users create custom C++ classes, subclassing appropriate abstract interfaces, in order to set up geometry, particle generation, custom data readout and physics modelling. At runtime, those classes are then instantiated and registered with the framework. Accordingly, the MCPL-**Geant4** integration takes the form of two such subclasses of **Geant4** abstract interfaces, which can be either directly used or

Listing 4: The G4MCPLGenerator class.

```

class G4MCPLGenerator : public G4VUserPrimaryGeneratorAction
{
public:

    G4MCPLGenerator(const G4String& inputFile);
    virtual ~G4MCPLGenerator();
    virtual void GeneratePrimaries(G4Event*);

protected:

    //Reimplement this to filter input particles (default
    //implementation accepts all particles):
    virtual bool UseParticle(const mcpl_particle_t*) const;

    //Reimplement this to change coordinates or weights of
    //input particles before using them (does nothing by
    //default):
    virtual void ModifyParticle(G4ThreeVector& position,
                                G4ThreeVector& direction,
                                G4ThreeVector& polarisation,
                                G4double& time,
                                G4double& weight) const;

private:
    // ..
};

```

further subclassed as needed: `G4MCPLGenerator` and `G4MCPLWriter`. They are believed to be compatible with any recent version of `Geant4` and were explicitly
300 tested with versions 10.00.p03 and 10.02.p01.

First, the `G4MCPLGenerator`, the relevant parts of which are shown in Listing 4, is a `G4VUserPrimaryGeneratorAction` which must be provided with the path to an MCPL file when constructed and in the `GeneratePrimaries` method of which the particles in the input file are used to generate `Geant4` events with a
305 single primary particle in each. If the file runs out of particles before the `Geant4` simulation is ended for other reasons, the `G4MCPLGenerator` graciously requests the `G4RunManager` to abort the simulation. Thus, a convenient way in which to use the entire input file for simulation is to launch the simulation with a very high number of events requested, as is done in the example in Listing 5⁵.

⁵Unfortunately, due to a limitation in the `G4RunManager` interface, this number will be limited by the highest number representable with a `G4int`, which on most modern platforms is 2147483647.

Listing 5: Example showing how to load particles from an MCPL file into a **Geant4** simulation.

```
#include "G4MCPLGenerator.hh"
#include "G4RunManager.hh"
#include <limits>

//Not shown here: Code defining MyGeometry and MyPhysicsList.

int main( int argc, char** argv ) {

    G4RunManager runManager;
    runManager.SetUserInitialization(new MyGeometry);
    runManager.SetUserInitialization(new MyPhysicsList);
    runManager.SetUserAction(new G4MCPLGenerator("myfile.mcpl"));
    runManager.Initialize();
    runManager.BeamOn(std::numeric_limits<G4int>::max());

    return 0;
}
```

310 In case the user wishes to choose only certain particles from the input file for simulation, the **G4MCPLGenerator** class must be sub-classed and the **UseParticle** method reimplemented, returning **false** for particles which should be skipped. Likewise, if it is desired to perform coordinate transformations or reweighing before using the loaded particles, the **ModifyParticle** method must
315 be reimplemented.

The **G4MCPLWriter** class, the relevant parts of which are shown in Listing 6, is a **G4VSensitiveDetector** which in the default configuration “consumes” all particles which, during a simulation, enters any geometrical volume(s) to which it is attached by the user and stores them into the specified MCPL file. At the
320 same time it asks **Geant4** to end further simulation of those particles (“killing” them). This strategy of killing particles stored into the file was chosen as a sensible default behaviour, as it prevents potential double-counting in the scenarios where a particle (or its induced secondary particles) would otherwise be able to enter a given volume multiple times. If it is desired to modify this strategy, the
325 user must sub-class **G4MCPLWriter** and reimplement the **ProcessHits** method, using calls to **storePreStep**, **storePostStep** and **Kill**, as desired. For reference, code responsible for the default implementation is shown in Listing 7. Likewise, to add MCPL user-flags into the file, the **UserFlagsDescription** and **UserFlags** methods must simply be reimplemented - the description naturally

Listing 6: The G4MCPLWriter class.

```
class G4MCPLWriter : public G4VSensitiveDetector
{
public:

    //Basic interface:

    G4MCPLWriter( const G4String& outputFile,
                  const G4String& name = "G4MCPLWriter" );
    virtual ~G4MCPLWriter();

    void AddComment( const G4String& );
    void AddData( const G4String& data_key,
                  size_t data_length,
                  const char* data );
    void EnableDoublePrecision();
    void EnablePolarisation();

    //Optional reimplement this to change default
    // "store-and-kill at entry" strategy:

    virtual G4bool ProcessHits( G4Step * step,
                                G4TouchableHistory* );

    //Optional reimplement these to add MCPL userflags:

    virtual G4String UserFlagsDescription() const { return ""; }
    virtual uint32_t UserFlags(const G4Step*) const { return 0; }

protected:
    //Methods that can be used if reimplementing ProcessHits():
    void storePreStep(const G4Step *);
    void storePostStep(const G4Step *);
    void Kill(G4Step *);

private:
    // ...
};
```

330 ending up as a comment in the output file.

In Listing 8 is shown how the G4MCPLWriter would typically be configured and attached to logical volume(s) of the geometry.

3.2. MCNP interface

Most users of MCNP are currently employing one of three distinct flavours:
 335 MCNPX[15, 16], MCNP5[17]^[todo: missing text or text in need of edit here] or MCNP6[18]. In the most typical mode of working with any of these software packages, users edit and launch MCNP through the use of text-based configuration files (so-called *input decks*), in order to set up details of the simulation including geometry, particle

Listing 7: The default ProcessHits implementation in the G4MCPLWriter class.

```
G4bool G4MCPLWriter::ProcessHits(G4Step * step, G4TouchableHistory*)
{
    //Only consider particle steps originating at the boundary
    //of the monitored volume:
    if ( step->GetPreStepPoint()->GetStepStatus() != fGeomBoundary )
        return false;

    //Store the state at the beginning of the step, but avoid
    //particles taking their very first step (this would double-
    //count secondary particles created at the volume edge):
    if ( step->GetTrack()->GetCurrentStepNumber() > 1 )
        storePreStep(step);

    //Tell Geant4 to stop further tracking of the particle:
    Kill(step);
    return true;
}
```

Listing 8: Example showing how to produce an MCPL file from a Geant4 simulation.

```
//Provide output filename when creating G4MCPLWriter instance:
G4MCPLWriter * mcplwriter = new G4MCPLWriter("myoutput.mcpl");

//Optional calls which add meta-data or change flags:
mcplwriter->AddComment("Some useful description here");
mcplwriter->AddData( ... );
mcplwriter->EnableDoublePrecision();
mcplwriter->EnablePolarisation();

//Register with G4SDManager and on one or more logical
//volumes to activate:
G4SDManager::GetSDMpointer()->AddNewDetector(mcplwriter);
alogvol->SetSensitiveDetector(mcplwriter);
```

generation, and data extraction. The latter typically resulting in the creation
340 of data files, ready for subsequent analysis.

Although it would be conceivable to write in-process **Fortran**-compatible
MCPL hooks for MCNP, such an approach would require users to undertake some
sort of compilation and linking procedure. This would thus be likely to impose
a change in working mode for the majority of MCNP users, in addition to possibly
345 requiring a special license for source-level access to MCNP. Instead, the MCNP–MCPL
interface presented here exploits the existing MCNP capability to stop and subse-
quently restart simulations at a user-defined set of surfaces through the **Source
Surface Write/Read (SSW/SSR)** functionality. As the name suggests, the state
parameters of particles crossing a given surface is stored on disk in dedicated
350 files, with the intentional use as a surface source in subsequent simulations with
the same MCNP setup. Throughout the present text and the MCPL distribution,
these files will be referred to as **SSW** files. Presumably<sup>[todo: missing text or text in need of
edit here]</sup>, the **SSW** file format is intended for this internal intermediate usage only,
it differs between different flavours of MCNP, and little effort has been made to
355 document it in publicly available manuals. Despite these obstacles, the **SSW**
format is stable enough that several existing MCNP-aware tools (e.g. [19, 20]<sup>[todo:
missing text or text in need of edit here]</sup>) have chosen to provide converters for this format,
with various levels of functionality, and it was thus deemed suitable also for the
needs of the MCPL project.

360 Thus, the MCPL distribution presented here includes dependency-free C code
for two standalone executables, **mcp12ssw** and **ssw2mcp1**, which users can invoke
from the command-line in order to convert between MCPL and **SSW** files. These
two commands will be discussed here, while users are referred to the relevant
MCNP manuals for details of how to set up their input decks to enable **SSW** input
365 or output in their MCNP simulations.^[todo: missing text or text in need of edit here] Note that
through usage of **ssw2mcp1** and **mcp12ssw**, it is even possible to transfer particles
between different flavours and versions of MCNP, which is otherwise not possible
with **SSW** files.

First, the **ssw2mcp1** command, for which the full usage instructions are

Listing 9: Usage instructions for the `ssw2mcpl` command.

```
Usage:
    ssw2mcpl [options] input.ssw [output.mcpl]

Converts the Monte Carlo particles in the input.ssw file (MCNP Surface
Source Write format) to MCPL format and stores in the designated output
file (defaults to "output.mcpl").

Options:
    -h, --help      : Show this usage information.
    -d, --double    : Enable double-precision storage of floating point values.
    -s, --surf      : Store SSW surface IDs in the MCPL userflags.
    -n, --nogzip    : Do not attempt to gzip output file.
    -c FILE         : Embed entire configuration FILE (the input deck)
                     used to produce input.ssw in the MCPL header.
```

370 shown in Listing 9, is in its most base invocation straight-forward to use. Simply provide it with the name of an existing SSW file to run on, and it will result in the creation of a new (compressed) MCPL file, `output.mcpl.gz`, containing a copy of all particles found in the SSW file. The MCNP flavour responsible for creating the SSW file is automatically detected, the resulting differences in the
375 file format are taken into account behind the scenes, and the detected MCNP version is documented as a comment in the header of the resulting MCPL file.

The only piece of information which is by default not transferred from the SSW particle state into the MCPL file is the numerical ID of the surface where the particle was captured in the MCNP simulation. By supplying the `-s` option,
380 `ssw2mcpl` will transfer those to the MCPL userflags field, and document this in the MCPL header. Additionally, while floating point numbers in the SSW file are always stored in double-precision, the transfer to MCPL will by default convert them to single-precision. This was chosen as the default behaviour to keep storage requirements low, as single-precision is often more than sufficient for
385 most studies. By supplying the `-d` option, `ssw2mcpl` will keep the numbers in double-precision in the MCPL file as well. Depending on compression and the applied flags, the resulting MCPL file will typically take up between 20–80% of the storage of the SSW file from which it was converted.

Finally it is possible, via the `-c FILE` flag, to point the `ssw2mcpl` command
390 to the input deck file used when producing the provided SSW file. Doing so will

result in a complete copy of that file being stored in the MCPL header as a binary data blob under the string key "`mcnp_input_deck`", thus providing users with a convenient snapshot in the MCPL file of the MCNP setup used. Unfortunately, it was not possible to automatise this procedure completely, and it thus relies on

395 the user to provide the correct input deck for a given SSW file (but the `ssw2mcpl` command does try to at least make sure the specified file is a text-file which contains at least the same problem title as the one found in the SSW file). The input deck embedded in a given MCPL file can later be inspected from the command line by invoking the command "`mcpltool -bmcnp_input_deck <file.mcpl>`".

400 Usage of the `mcpl2ssw` command, for which the full usage instructions are shown in Listing 10, is slightly more involved: in addition to an input MCPL file, the user must also supply a reference SSW file in a format suitable for the MCNP setup in which the file is to be used as input. The need for this added complexity stems from the constraint that the SSW format is merely intended

405 as an internal format in which it is possible to stop and restart particles while remaining within a given setup of an MCNP simulation (including at the very least the MCNP version and the configuration of the geometrical surfaces involved in the **Source Surface Write/Read** procedure). Thus, for maximal robustness, the user must supply a reference SSW file which was produced by the setup in

410 which the SSW file created with `mcpl2ssw` is to be used (it does not matter how many particles the reference file contains). What will actually happen, is that in addition to the particle state data itself, the newly created SSW file will contain the exact same header as the one in the reference SSW file, apart from the fields related to the number of particles in the file.

415 Additionally, the user must consider carefully which MCNP surface IDs the particles from the MCPL file should be associated with, once transfered to the SSW file. By default it will assume that the MCPL `userflags` field contains exactly this ID, but more often than not, users will have to specify a global surface ID for all of the particles through the `-s<ID>` command-line option for the `mcpl2ssw`

420 command.

Finally, note that SSW files do not contain polarisation information, and any

Listing 10: Usage instructions for the `mcpl2ssw` command.

```
Usage:

    mcpl2ssw [options] <input.mcpl> <reference.ssw> [output.ssw]

Converts the Monte Carlo particles in the input MCPL file to SSW format
(MCNP Surface Source Write) and stores the result in the designated output
file (defaults to "output.ssw").

In order to do so and get the details of the SSW format correct, the user
must also provide a reference SSW file from the same approximate setup
(MCNP version, input deck...) where the new SSW file is to be used. The
reference SSW file can of course be very small, as only the file header is
important (the new file essentially gets a copy of the header found in the
reference file, except for certain fields related to number of particles
whose values are changed).

Finally, one must pay attention to the Surface ID assigned to the
particles in the resulting SSW file: Either the user specifies a global
one with -s<ID>, or it is assumed that the MCPL userflags field in the
input file is actually intended to become the Surface ID. Note that not
all MCPL files have userflag fields and that valid Surface IDs are
integers in the range 1-999999.

Options:

    -h, --help      : Show this usage information.
    -s<ID>          : All particles in the SSW file will get this surface ID.
    -l<LIMIT>       : Limit the number of particles transferred to the SSW file
                      (defaults to 2147483647, the maximal SSW capacity).
```

polarisation info in the input MCPL file will consequently be discarded in the translation. Likewise, in cases where the input MCPL file contain one or more particles whose type does not have a representation in the relevant flavour of MCNP, they will be ignored with suitable warnings.

3.3. *McStas* and *McXtrace* interfaces

Recent releases of the neutron ray tracing software package *McStas* [21, 22] (version 2.3 and later) and its X-ray sibling package *McXtrace* [23] (version 1.3 and later) both include MCPL-interfaces. Although *McStas* and *McXtrace* are two distinct software packages, they are implemented upon a common technological platform, *McCode*, and the discussions here will for simplicity use the term *McCode* where the instructions are otherwise identical for users of the two packages.

The particle model adopted in *McCode* is directly compatible with MCPL. In essence, apart from mere unit conversions, particles are read from or written to MCPL files at one or more predefined logical points defined in the *McCode* configuration files (the *instrument files*). Specifically, two new components, *MCPL_input*

Listing 11: Code enabling MCPL input in its simplest form.

```
COMPONENT vin = MCPL_input( filename="myfile.mcpl" )  
AT(0,0,0) RELATIVE Origin
```

Listing 12: Code enabling MCPL input, selecting particles in a given energy range.

```
COMPONENT vin = MCPL_input( filename="myfile.mcpl",  
                             Emin=12, Emax=100 )  
AT(0,0,0) RELATIVE Origin
```

and `MCPL_output`, have been provided, which users can enable by adding entries at relevant points in their instrument files as is usual when working with `McCode`.

First, when using the `MCPL_input` component, particles are directly read from an MCPL input file and injected into the simulation at the desired point, thus playing the role of a source model. In Listing 11 is shown how, in its simplest form, users would insert an `MCPL_input` component in their instrument file. This will result in the MCPL file being read in its entirety, and all found neutrons (for `McStas`) or gamma particles (for `McXtrace`) are traced through the `McCode` simulation.^[todo: missing text or text in need of edit here] In Listing 12 is indicated how the user can additionally impose an allowed energy range when loading particles by supplying the `Emin` and `Emax` parameters. The units are meV and keV respectively for `McStas` and `McXtrace`. Thus, the code in Listing 12 would select 12–100meV neutrons in `McStas` and 12-100keV gammas when used in `McXtrace`.

For technical reasons, the number of particles to be simulated in `McCode` must be fixed at initialisation time. Thus, the number of particles will be set to the total number of particles in the input file, as this is provided through the corresponding MCPL header field. If and when a particle is encountered which can not be used (due to having a wrong particle type or energy), it will lead to an empty event in which no particles leave the source. At the end of the run, the number of particles skipped over will be summarised for the user. This approach obviates the need for running twice over the input file, as well as avoiding the

Listing 13: Code enabling MCPL output in its simplest form.

```
COMPONENT mcplout = MCPL_output( filename="myoutput.mcpl" )  
AT(0,0,0) RELATIVE PREVIOUS
```

potential introduction of statistical bias from reading a partial file.

Note that if running `McCode` in parallel processing mode using `MPI` [todo: missing text or text in need of edit here], each process will operate on all particles in the entire file, but the particles will get their statistical weights reduced accordingly upon
465 load. This behaviour is not specific to the `MCPL_input` component, but is a general feature of how multiprocessing is implemented in `McCode`.

This is probably one of the most likely use-cases of the `McStas`/`MCPL` interface where a neutron moderator is modelled using e.g. `MCNP` and the rest of the instrument is modelled using `McStas`. [todo: missing text or text in need of edit here]

470 When adding an `MCPL_output` component to an `McCode` instrument file, the complete state of all particles reaching that component are written to the requested output file. In Listing 13 is shown how, in its simplest form, users would insert such a component in their instrument file, and get particles written with coordinates relative to the component preceding it, into the output file. For
475 reference, a copy of the complete instrument file is stored in the `MCPL` header as a binary data blob under the string key [todo: missing text or text in need of edit here]. This feature provides users with a convenient snapshot of the generating setup. The instrument file contained in a given `MCPL` can be inspected from the command line by invoking the command “`mcpltool -b<key> <file.mcpl>`”.

480 If running `McCode` in parallel processing mode using `MPI`, each process will in this case create a separate output file named `myoutput-idx.mcpl` where `idx` is the process number (assuming `filename="myoutput.mcpl"` as in Listing 13). The resulting files can subsequently be merged using the `mcpltool` command if desired. It is envisioned that a future version of the `MCPL_output` component
485 will automatically handle this merging for the user.

To avoid generating unnecessarily large files, the `MCPL_output` component stores particle state data using the global PDG code feature (cf. section 2.1),

Listing 14: Code enabling MCPL output with polarisation and double precision numbers.

```
COMPONENT mcplout = MCPL_output( filename="myoutput.mcpl",
                                polarisationuse=1,
                                doubleprec=1 )
AT(0,0,0) RELATIVE PREVIOUS
```

Listing 15: Code enabling MCPL output with custom user-flags information.

```
/* some upstream component setting a variable (customvar) */
COMPONENT some_comp = Some_Component( /* some parameters here */ )
AT(0,0,0) ABSOLUTE
EXTEND %{
    customvar=(uint32_t) mcget_run_num();
%}

/* MCPL output capturing customvar into MCPL user-flags */
COMPONENT vout = MCPL_output( filename="myoutput.mcpl",
                              userflag=customvar,
                              userflagcomment="Particle Id" )
AT(0,0,0) RELATIVE PREVIOUS
,
```

uses single precision floating point numbers, and does *not* by default store polarisation vectors. The two latter settings may be changed by the user through the `polarisationuse` and `doubleprec` parameters respectively, as shown in listing 14.

Finally, if desired, custom information might be stored per-particle into the MCPL user-flags field for later reference. This could be any property, such as for instance the number of reflections along a neutron guide, or the type of scattering process in a crystal, etc. Listing 15 shows a simple example of this where the particle ID, in the form of its McCode ray number (returned from the McCode library function `mcget_run_num`), is stored into the user-flags field. A string, `userflagcomment`, is required in order to describe the significance of the extra data, and will end up as a comment in the resulting MCPL file.

4. Example uses

[todo: missing text or text in need of edit here]

5. Conclusion and outlook

[todo: missing text or text in need of edit here]

505 Latest version of code and documentation is available at [24].

Acknowledgements

This work was supported in part by the EU Horizon 2020 framework, Bright-
nESS project 676548. [todo: missing text or text in need of edit here]

References

- 510 [1] T. Sjöstrand, et al., A Proposed Standard Event Record, in: G. Altarelli,
et al. (Eds.), Z physics at LEP 1, Vol. 3, CERN, Geneva, 1989, Ch. 6.2,
pp. 327–330. doi:10.5170/CERN-1989-008-V-3.
- [2] M. Dobbs, J. B. Hansen, The HepMC C++ Monte Carlo event record for
High Energy Physics, Computer Physics Communications 134 (1) (2001)
515 41 – 46. doi:10.1016/S0010-4655(00)00189-2.
- [3] J. Alwall, et al., A standard format for Les Houches Event Files, Computer
Physics Communications 176 (4) (2007) 300 – 304. doi:10.1016/j.cpc.
2006.11.010.
- [4] K A Olive, et al. (Particle Data Group), Review of Particle Physics, Chinese
520 Physics C 38 (9) (2014) 090001. doi:10.1088/1674-1137/38/9/090001.
- [5] T. Engelhardt, C. Dachsbacher, Octahedron environment maps, Proceed-
ings of the Vision, Modeling, and Visualization Conference 2008, VMV
2008, Konstanz, Germany (2008) 383–388.
- [6] Z. H. Cigolle, S. Donow, D. Evangelakos, M. Mara, M. McGuire, Q. Meyer,
525 A survey of efficient representations for independent unit vectors, Journal
of Computer Graphics Techniques (JCGT) 3 (2) (2014) 1–30.
URL <http://jcgt.org/published/0003/02/01/>

- [7] S. Agostinelli, et al., GEANT4: A Simulation toolkit, Nucl. Instrum. Meth. A506 (2003) 250–303. doi:10.1016/S0168-9002(03)01368-8.
- 530 [8] J. Allison, et al., Geant4 developments and applications, IEEE Trans. Nucl. Sci. 53 (2006) 270. doi:10.1109/TNS.2006.869826.
- [9] T. Kittelmann, et al., Geant4 based simulations for novel neutron detector development, J. Phys: Conf. Ser. 513 (2014) 022017. doi:10.1088/1742-6596/513/2/022017.
- 535 [10] L. P. Deutsch, GZIP file format specification version 4.3, RFC 1952 (May 1996). doi:10.17487/RFC1952.
- [11] L. P. Deutsch, DEFLATE Compressed Data Format Specification version 1.3, RFC 1951 (May 1996). doi:10.17487/RFC1951.
- [12] Jean-Loup Gailly and Mark Adler, ZLIB Compression library, v1.2.8.
540 URL <http://zlib.net/>
- [13] L. P. Deutsch, J.-L. Gailly, ZLIB Compressed Data Format Specification version 3.3, RFC 1950 (May 1996). doi:10.17487/RFC1950.
- [14] K. Martin, B. Hoffman, Mastering CMake: A Cross-Platform Build System, Kitware Inc, 2015.
- 545 [15] L. S. Waters, et al., The MCNPX Monte Carlo radiation transport code, AIP Conf. Proc. 896 (2007) 81–90, [,81(2007)]. doi:10.1063/1.2720459.
- [16] D. B. Pelowitz, et al., MCNPX 2.7.0 Extensions, Tech. Rep. LA-UR-11-02295, Los Alamos National Laboratory (2011).
- 550 [17] X.-. M. C. Team, MCNP – A General Monte Carlo N-Particle Transport Code, Version 5, Tech. Rep. LA-UR-03-1987, Los Alamos National Laboratory (2003 (revised 2008)).
- [18] D. B. Pelowitz, et al., MCNP6 Users Manual, , Tech. Rep. LA-CP-13-00634, Los Alamos National Laboratory (2013).

- [19] the PyNE Development Team, PyNE: The Nuclear Engineering Toolkit
555 (2014).
URL <http://pyne.io>
- [20] E. Klinkby, et al., Interfacing MCNPX and McStas for simulation of
neutron transport, Nucl. Instrum. Meth. A700 (2013) 106 – 110. doi:
10.1016/j.nima.2012.10.052.
- 560 [21] K. Lefmann, K. Nielsen, Mcstas, a general software package for neutron
ray-tracing simulations, Neutron News 10 (3) (1999) 20–23. doi:10.1080/
10448639908233684.
- [22] P. Willendrup, E. Farhi, K. Lefmann, McStas 1.7 - a new version of the
flexible Monte Carlo neutron scattering package, Physica B: Condensed
565 Matter 350 (1âĂŞ3, Supplement) (2004) E735 – E737. doi:10.1016/j.
physb.2004.03.193.
- [23] E. Bergbäck Knudsen, et al., McXtrace, Journal of Applied Crystallography
46 (3) (2013) 679–696. doi:10.1107/S0021889813007991.
- [24] T Kittelmann and others, Website of MCPL: Monte Carlo Particle Lists
570 (2016).
URL <https://mctools.github.io/mcpl/>